

Boichitro-MoE: consolidated results and Q1 evaluation

Snapshot: 2026-07-22 06:12 KST

Report status: interim, evidence-audited snapshot

Pipeline status at snapshot: RUNNING

Submission status: NOT READY

This report inventories every finalized result family currently present in the workspace. Development, previously accessed test, smoke-test, and missing locked evidence are kept separate. Values from unfinished checkpoints are not treated as scientific results.

Executive evaluation

The project is methodologically ambitious and unusually strong in provenance, leakage control, artifact tracking, and reproducibility. However, its main scientific claim is not yet testable: two M2 and all three M3 main runs, protocol freezing, locked evaluation, confirmatory statistics, robustness, routing analysis, and human evaluation are unfinished.

The raw architecture results remain mixed, but the previously weak system-level evaluation has now been repaired under a source-blind contract:

- M1 improves over M0 by **+0.863 macro chrF++** on normalization validation, but loses **0.661 worst-dialect chrF++**, increases replay degradation by **0.818 percentage points**, and has worse selection-time calibration.
- M1 improves M0 identification regional macro-F1 by only **+0.0053**.
- The earlier **50.462** word-rewrite score used gold dialect metadata. It is now correctly labeled as an oracle, not a deployable comparator.
- The corrected inferred-supported-dialect rewrite scores **50.005 macro chrF++** and **30.388 worst-dialect chrF++** with no gold metadata.
- A semantic-group-held-out, five-fold, reference-free candidate selector reaches **52.667 macro chrF++** and **34.851 worst-dialect chrF++** over six M0/M1 runs. Its paired gain over the fair rewrite is **+2.662** (95% dependence-preserving hierarchical bootstrap CI **[2.086, 3.268]**; paired randomization **p=0.000200**).
- The frozen neural/SVM identification blend reaches **0.8009 regional macro-F1**, a paired **+0.0390** over raw neural predictions (95% CI **[0.0309, 0.0479]**; **p=0.000200**), with **0.0440 ECE-15**.

These fusion values are development-only and selection-conditioned: the same development rows informed the selected family/threshold or blend weight. Their intervals preserve shared semantic-group dependence across runs and dialects, but do not include selection uncertainty and are exploratory, not confirmatory. Raw neural outputs remain primary for architecture and source-OOD inference, and the fixed fusion is saved as a separate system view for the later locked evaluation.

Accordingly, the evidence presently supports the engineering contribution and some modest dense-to-Switch development gains. It does **not** support Boichitro-MoE superiority, source-OOD robustness, state-of-the-art performance, or publication readiness.

Current Q1 readiness rating: 5.5/10. The evaluation design and deployable development systems are materially stronger, but submission now would still be premature because the central M3-vs-M2 locked claim and required human evidence do not yet exist.

Evidence status

Result family	Finalized status	Admissible interpretation
Data engineering and leakage gates	Pass internally	Training-ready; native linguistic validation pending
Tokenizer selection	Complete	Validation-only, three seeds, no test access
Dense foundation and three continuations	Complete	Descriptive, one foundation seed
Development pilots	Complete	Hyperparameter and topology selection only
M0 main task runs	3/3	Development validation only
M1 main task runs	3/3	Development validation only
Fair source-blind fusion	Complete on M0/M1	Validation-only; fixed before locked evaluation
M2 main task runs	1/3 finalized	Seed 1701 is complete; insufficient for an aggregate claim
M3 main task runs	0/3	No result
Fixed fusion transfer to M2/M3	1/6	M2 seed 1701 diagnostic only; shared validation rows
Locked main neural evaluation	0/12	No confirmatory result
Confirmatory and optimization ablations	0/30	No causal component evidence
Human dataset review	0/230	No native-quality corpus claim
Blinded human system ratings	Missing	No adequacy/fluency claim

The readiness audit passes **10 of 29** gates. The 19 failed gates are blocking, not cosmetic.

Dataset results

Final admitted data

Quantity	Rows
Authentic normalization pairs	54,598
Traceable synthetic normalization pairs, train only	3,325
All normalization rows	57,923

Quantity	Rows
Identification rows	122,353
Real romanized source-OOD challenge	1,342
Unique tokenizer-training texts	100,236
Recorded exclusion decisions	137,386
Quarantined legacy-derived rows	51,101

Normalization contains 39,622 training, 4,289 validation, 5,033 IID-test, and 8,979 OOD rows. The OOD total consists of 7,592 supported-dialect source-OOD rows plus 1,387 RAJ zero-shot rows. Identification contains 92,897 training, 10,238 validation, 11,264 IID-test, 6,803 source-OOD, and 1,151 external rows.

The engineering gate passes ten automatic checks: the 13-label inventory is consistent, exact/compact train-evaluation overlap is absent, cross-label conflicts are removed, synthetic data are train-only and excluded from ID, locked OOD sources never train, the legacy archive is absent, and licenses are resolved. The sole failed publication gate is native review: **0 of 230 rows** have been reviewed.

Dataset limitations

- Normalization training is concentrated in CHI (36.76%) and SYL (36.54%); together they account for 73.3% of training rows.
- RAJ has no normalization training or validation pair, so its 1,387-row score must be labelled zero-shot and cannot be merged into supported-dialect OOD.
- KHU, NAR, and RAN lack defensible independent-source normalization tests.
- The dataset is internally verified but cannot yet be described as native-speaker validated or cleared for public redistribution.

Tokenizer results

Twelve intrinsic candidates passed the mechanical gates; four entered the three-seed, 2M-token-per-seed proxy comparison.

Candidate	Proxy BPC mean $\pm$ SD $\downarrow$	Worst-dialect BPC $\downarrow$	Tokens/character $\downarrow$	Tokens/s $\uparrow$	Selected
byte-BPE balanced 16k	3.3807 $\pm$ 0.0113	3.5934	0.6859	153,304	No
Unigram balanced 32k	2.7482 $\pm$ 0.0114	2.9518	0.2383	72,811	No
WordPiece balanced 32k	2.6675 $\pm$ 0.0023	2.8898	0.2268	71,746	No
WordPiece natural 32k	2.6600 $\pm$ 0.0088	2.8753	0.2245	74,345	Yes

The frozen choice is well supported within the screened set: it has the best mean and worst-dialect proxy BPC and the lowest token cost among the shortlisted non-byte tokenizers. Selection used no test data. The limitation is breadth: the proxy study establishes the best local candidate, not superiority to every modern Bangla tokenizer.

Pretraining corpus and foundation results

The revision-pinned FineWeb-2 Bangla build saw 585,619 documents, accepted 567,746, and produced 300,004,991 training tokens plus 3,000,629 validation tokens. Recorded rejections include 3,505 low-Bangla-ratio documents, 2,476 direct benchmark matches, 49 compact benchmark matches, 130 overlong documents, and 69 invalid-Unicode documents. All five source parquet hashes were verified.

Fixed-budget foundation comparison

Run	Tokens	Final validation BPC ↓	Change from 300M dense foundation	Training tok/s
F dense foundation	300.0M	1.230329	—	23,853
M0 dense continuation	200.0M	1.203249	−2.201%	16,877
M1 Switch continuation	200.0M	1.199294	−2.523%	14,100
M2 shared top-2 MoE continuation	200.0M	1.202603	−2.254%	18,223

M1 is best by validation BPC, but its advantage over M0 is only 0.00396 BPC (0.329%). M2 improves on M0 by 0.00065 BPC (0.054%). These are single-seed foundation results, so they are descriptive and cannot establish a stable architecture ranking.

Router diagnostics passed the non-finite and persistent-collapse guards. M1's mean/final load CV is 0.223/0.258; M2's is 0.147/0.043. One isolated M1 CV exceedance and no post-stability M2 exceedance were recorded.

Development-pilot results

Mature-checkpoint continuation learning rate

Muon / AdamW LR	Final BPC ↓	Relative change	Eligible Decision
0.001 / 0.000015	1.226813	−0.286%	Yes Selected
0.002 / 0.000030	1.232059	+0.141%	No Rejected
0.005 / 0.000075	1.254107	+1.933%	No Rejected
0.010 / 0.000150	1.293376	+5.124%	No Rejected

The monotonic deterioration at larger restart rates is credible evidence that a mature checkpoint required a low continuation LR.

Dense-to-MoE recovery

Strategy	Initial BPC	Final BPC ↓	Maximum regression	Eligible
Abrupt bank release	1.230335	1.262179	+3.678%	No
Unbanked transfer	1.397035	1.255871	0.000%	Negative control
Random initialization	3.519622	2.559432	0.000%	Negative control
Annealed cross-bank	1.230335	1.264109	+2.745%	No
Permanent paired bank	1.230335	1.228724	0.000%	Selected

Only permanent paired-bank routing satisfied the preregistered transient and final regression guards. This is one of the clearest positive pilot findings.

Switch-router repair

Router strategy	Final BPC ↓	Maximum load CV ↓	Final load CV ↓	Eligible
Loss-free straight-through	1.227010	0.5061	0.5061	No
Auxiliary-balance straight-through	1.226940	0.2617	0.0675	Yes

The auxiliary-balance router is both marginally better in BPC and substantially more balanced, justifying its selection.

Stage-S retention schedule

Candidate	Norm/replay fraction	Macro chrF++ ↑	Worst dialect ↑	Replay degradation ↓	Selected
Rejected default	0.55 / 0.10	44.714	34.214	15.910%	No
ret25 balanced	0.40 / 0.25	41.158	31.917	1.336%	No
ret25 conservative	0.40 / 0.25	41.072	30.621	0.831%	No
ret35 balanced	0.30 / 0.35	41.186	31.910	0.972%	Yes
ret35 conservative	0.30 / 0.35	40.745	31.285	0.362%	No

The selected schedule correctly respects the 5% replay guard, but the trade-off is large: it gives up 3.53 chrF++ relative to the rejected high-forgetting schedule. The choice is protocol-consistent, not evidence that the selected score is competitive.

Main task development results

M0 and M1 have finalized three-seed development runs; M2 seed 1701 is also finalized while its other

seeds and all M3 runs remain in progress. All values below are validation-selection results and must not be reported as locked test results.

Per-seed results

Model	Seed	Norm macro chrF++ ↑	Worst dialect ↑	Replay degradation ↓	ID regional F1 ↑	ID macro-F1 ↑	Accuracy ↑	ECE ↓
M0	1701	41.186	31.910	0.972%	0.7605	0.7726	0.8436	0.0900
M0	2903	41.232	32.184	1.128%	0.7519	0.7648	0.8426	0.0882
M0	4307	41.035	32.733	1.144%	0.7652	0.7769	0.8497	0.1131
M1	1701	42.278	31.534	1.898%	0.7670	0.7788	0.8533	0.1175
M1	2903	41.802	31.504	1.705%	0.7678	0.7796	0.8498	0.1197
M1	4307	41.963	31.806	2.094%	0.7586	0.7706	0.8456	0.1062
M2	1701	41.728	31.363	1.168%	0.7533	0.7657	0.8404	0.0965

Three-seed aggregates and paired M1–M0 effects

Metric	M0 mean ± SD	M1 mean ± SD	Paired difference
Normalization macro chrF++ ↑	41.151 ± 0.103	42.014 ± 0.242	+0.863 ± 0.267
Worst-dialect chrF++ ↑	32.276 ± 0.419	31.615 ± 0.167	−0.661 ± 0.276
Replay degradation ↓	1.081 ± 0.095%	1.899 ± 0.195%	+0.818 ± 0.209 pp
ID accuracy ↑	0.8453 ± 0.0038	0.8495 ± 0.0039	+0.0042 ± 0.0073
ID balanced accuracy ↑	0.7582 ± 0.0010	0.7684 ± 0.0031	+0.0102 ± 0.0038
ID macro-F1 (13) ↑	0.7714 ± 0.0062	0.7763 ± 0.0050	+0.0049 ± 0.0107
ID regional macro-F1 ↑	0.7592 ± 0.0067	0.7645 ± 0.0051	+0.0053 ± 0.0113
ID ECE ↓	0.0971 ± 0.0138	0.1145 ± 0.0073	+0.0174 ± 0.0211
ID Brier ↓	0.2496 ± 0.0062	0.2589 ± 0.0055	+0.0093 ± 0.0117

With only three seeds and development reuse, these differences are descriptive. The paired signs show a real trade-off: M1 improves mean task scores but is worse on the normalization tail, retention, ECE, and Brier score.

Normalization validation by dialect

Dialect	Word rewrite	M0 mean $\pm$ SD	M1 mean $\pm$ SD	M1–rewrite
BAR	50.978	42.634 $\pm$ 1.424	42.841 $\pm$ 0.803	−8.137
CHI	62.192	50.019 $\pm$ 1.033	50.077 $\pm$ 0.358	−12.115
KHU	34.604	34.057 $\pm$ 1.183	36.118 $\pm$ 1.102	+1.513
MYM	77.089	52.674 $\pm$ 1.181	53.783 $\pm$ 0.751	−23.306
NAR	41.184	35.703 $\pm$ 1.033	38.852 $\pm$ 2.032	−2.332
NOA	51.727	38.058 $\pm$ 0.742	38.772 $\pm$ 0.532	−12.955
RAN	30.575	32.311 $\pm$ 0.478	31.615 $\pm$ 0.167	+1.039
SYL	55.349	43.753 $\pm$ 0.107	44.056 $\pm$ 0.292	−11.293

M1 beats the word-rewrite control only on KHU and RAN. The largest deficit is MYM (−23.31 chrF++). This is the most serious current performance finding and must not be hidden by an architecture-only comparison.

Identification validation by class

Label	Support	M0 F1 mean $\pm$ SD	M1 F1 mean $\pm$ SD	M1–M0
BAR	413	0.8365 $\pm$ 0.0150	0.8451 $\pm$ 0.0098	+0.0086
CHI	1,991	0.9446 $\pm$ 0.0005	0.9481 $\pm$ 0.0031	+0.0035
KHU	86	0.4613 $\pm$ 0.0779	0.4613 $\pm$ 0.0442	0.0000
KIS	797	0.7559 $\pm$ 0.0059	0.7632 $\pm$ 0.0085	+0.0073
MYM	302	0.7826 $\pm$ 0.0049	0.7559 $\pm$ 0.0080	−0.0267
NAR	865	0.7803 $\pm$ 0.0135	0.7842 $\pm$ 0.0110	+0.0039
NOA	260	0.7836 $\pm$ 0.0109	0.7853 $\pm$ 0.0131	+0.0017
NSD	576	0.6791 $\pm$ 0.0189	0.7015 $\pm$ 0.0172	+0.0224
RAJ	91	0.7634 $\pm$ 0.0470	0.7696 $\pm$ 0.0159	+0.0062
RAN	680	0.7670 $\pm$ 0.0070	0.7825 $\pm$ 0.0107	+0.0155
STD	2,251	0.9185 $\pm$ 0.0006	0.9186 $\pm$ 0.0042	+0.0001
SYL	1,407	0.8997 $\pm$ 0.0030	0.9058 $\pm$ 0.0045	+0.0061

Label Support M0 F1 mean $\pm$ SD M1 F1 mean $\pm$ SD M1-M0

TAN 519 0.6561 $\pm$ 0.0070 0.6713 $\pm$ 0.0191 +0.0152

KHU remains the weakest class at 0.461 F1, consistent with its very small training and validation support. M1 improves most classes but materially harms MYM.

Source-blind evaluation repair (development only)

System	Primary metric	Worst-group metric	Paired gain	95% CI
Fair supported-dialect rewrite	50.005 macro chrF++	30.388 chrF++	reference	—
Normalization selector V2, six-run semantic-group OOF mean	52.667 macro chrF++	34.851 chrF++	+2.662	[2.086, 3.268]
Raw neural ID, six-run mean	0.7618 regional macro-F1	—	reference	—
Neural/SVM ID blend	0.8009 regional macro-F1	0.4508 F1	+0.0390	[0.0309, 0.0479]

The normalization selector uses only source text, its neural and rewrite candidates, and dialect probabilities inferred from source text. References, gold dialects, source IDs, and evaluation-track labels are forbidden inference features. All rows and six repeated outputs sharing a semantic group remain in the same fold. Resampling assigns one multiplicity or swap to each global semantic group and reuses it across runs and cross-dialect realizations. Both selection-conditioned paired randomization diagnostics give two-sided $p=0.000200$ with 5,000 draws. These p -values are exploratory because fusion settings reused development data. The selected artifacts and their SHA-256 hashes are registered for protocol freeze; neither pilot accessed a test split. An aggregate-only cross-task firewall additionally found zero exact, compact, or registered SimHash/Jaccard near-duplicate overlap between all normalization evaluation inputs and the 92,897 inputs used to train the dialect classifier; it emitted no protected text and was not used for model selection.

The fixed artifacts have also transferred without retuning to the first later architecture run (1/6 planned M2/M3 runs). On M2 seed 1701, normalization rises from **41.728 raw** to **54.858 fused** macro chrF++ (worst dialect **36.378**), and identification regional macro-F1 rises from **0.7533 raw** to **0.8046 fused** with **0.0489 ECE-15**. This is a useful architecture-transfer diagnostic, but it reuses the same validation rows and one seed cannot establish stability or independent confirmation.

Fixed classical baselines

These baselines were evaluated before the neural protocol freeze. They are useful descriptive floors, but the test sets are not pristine. This limitation is permanently disclosed and cannot be repaired by deleting outputs.

Normalization

Model / track	Rows	Macro chrF++ ↑	Worst dialect ↑	Macro BLEU ↑	TER ↓	CER ↓	Exact match ↑
Copy / validation	4,289	38.401	23.764	10.021	63.106	0.348	0.0415
Inferred supported-dialect rewrite / validation	4,289	50.005	30.388	—	—	—	—
Gold-dialect rewrite oracle / validation	4,289	50.462	30.575	23.774	45.079	0.258	0.0874
Copy / IID test	5,033	39.260	25.853	10.750	62.102	0.334	0.0296
Gold-dialect rewrite oracle / IID test	5,033	50.916	33.423	22.912	44.316	0.250	0.0757
Copy / source OOD	7,592	38.652	25.319	8.602	70.984	0.409	0.1616
Gold-dialect rewrite oracle / source OOD	7,592	44.607	31.767	15.191	63.849	0.378	0.1927
Copy / RAJ zero-shot	1,387	48.642	48.642	4.912	59.002	0.357	0.2884
Gold-dialect rewrite oracle / RAJ zero-shot	1,387	48.642	48.642	4.912	59.002	0.357	0.2884

The legacy rewrite equals copy on RAJ because RAJ has no learned rewrite inventory. All legacy test rewrite values above used gold dialect and are diagnostic oracle values. The fair inferred-dialect baseline is evaluated on locked tests only after protocol freeze.

Identification

Model / track	Accuracy ↑	Macro-F1 (13) ↑	Regional macro-F1 ↑	ECE ↓	Worst present F1 ↑
SVM / validation	0.8550	0.7838	0.7734	0.0113	0.4138
SVM / IID test	0.8540	0.7783	0.7669	0.0111	0.4118
SVM / source OOD	0.4724	0.2229	0.3738	0.1891	0.0343
SVM / external transcript	0.9705	0.5943	0.9658	0.0146	0.8739
SGD, 5-seed mean / validation	0.8510	0.7768	0.7664	0.0566	0.3850
SGD, 5-seed mean / IID test	0.8498	0.7759	0.7653	0.0534	0.4377
SGD, 5-seed mean / source OOD	0.4686	0.2168	0.3630	0.1624	0.0257

Model / track	Accuracy ↑	Macro-F1 (13) ↑	Regional macro-F1 ↑	ECE ↓	Worst present F1 ↑
SGD, 5-seed mean / external	0.9672	0.5917	0.9614	0.0415	0.8483

The SVM's source-OOD collapse (regional macro-F1 0.374 versus 0.773 on validation) strongly motivates source-robust modeling. Conversely, its high IID score and excellent calibration mean the neural models must beat a strong, cheap baseline—not merely the majority classifier.

The external macro-F1 (13) is low because absent labels are included in that 13-way scalar, while regional macro-F1 is computed over labels present in the external track. Both must be reported with the coverage definition.

Systems results on NVIDIA GB10

Model	Total params	Active params/token	Throughput tok/s ↑	Peak memory GiB ↓	Slowdown vs M0
M0 dense	83.79M	83.79M	15,337	2.83	1.00×
M1 Switch	381.11M	83.84M	5,211	5.69	2.94×
M2 standard MoE	168.77M	83.84M	8,272	3.84	1.85×
M3 Boichitro	168.79M	83.86M	8,224	3.84	1.86×

Active parameters are tightly matched, but wall-clock compute is not. M2/M3 run at only about 54% of dense throughput, and M1 at 34%. M3 adds just 0.58% throughput overhead and 0.0068 GiB over M2, which is favorable for the proposed mechanisms. Nevertheless, any “compute matched” wording must distinguish active parameter matching from measured time/energy matching.

Smoke-test-only output

A tiny M3 end-to-end smoke run evaluated deterministic prefixes of 8 normalization and 16 identification examples per track. It scores 2.111 IID chrF++, 0 source-OOD/romanized chrF++, and zero ID accuracy. Its manifest explicitly states `smoke_only_not_scientific_evidence`; these values demonstrate pipeline execution only and must never appear as model-quality results.

Software, tables, and figures

- Regression suite: **107 passed**, with one non-failing PyTorch GPU-capability warning.
- Journal figure suite: **44 PNG/PDF pairs**, exceeding the requested minimum.
- Figure validation: **398/398 checks passed** for format, hashes, resolution, captions, and source data.
- Evidence classes: **40 development-only** and **4 prior-test descriptive**; none is mislabeled as locked evidence.
- Paper tables: **10 reproducible table families**.
- Manuscript: 3,154 words, structurally sound but still a short journal-neutral draft with withheld results, abbreviated references, and required metadata.

The figure engineering is strong, but 44 technically valid figures do not replace locked evidence. A final

paper should curate a small main set and move most diagnostics to supplementary material.

Claim-by-claim evaluation

Proposed claim	Current verdict	Reason
Provenance-first, leakage-controlled benchmark	Provisionally supported	Automatic gates pass; native review and redistribution clearance remain
Frozen tokenizer chosen without test access	Supported	Three-seed proxy and immutable hashes exist
Stable dense-to-MoE upcycling	Supported developmentally	Paired-bank pilot passes guards; one pilot seed
M1 improves over dense M0	Mixed	Mean metrics improve slightly; tail performance, replay, and calibration worsen
M3 outperforms M2 on source-OOD normalization	Untested	One M2 development seed exists; no M3 run or locked manifest exists
Fixed source-blind system is competitive on validation	Supported developmentally	Selector reaches 52.667 chrF++; ID blend reaches 0.8009 F1 with paired positive exploratory intervals
Source-invariant routing works	Untested	Locked routing specialization analysis missing
Robust under perturbation/romanization	Untested	Registered robustness evaluation missing
Positive development fusion diagnostics	Supported, exploratory	Dependence-preserving intervals exclude zero, but selection reused these rows and locked confirmation is missing
Human-acceptable normalization	Untested	No blinded native ratings

Q1 scorecard

Dimension	Score	Assessment
Novelty and research framing	8/10	Important low-resource/source-confounding question
Data and protocol engineering	8/10	Strong controls; native and redistribution gates remain
Reproducibility and artifact discipline	9/10	Hashes, registries, manifests, tests, source data
Systems evidence	7/10	Good active-parameter accounting; no energy/time matching

Dimension	Score	Assessment
Current model competitiveness	6/10	Raw neural normalization trails, but fair fixed fusion now exceeds the source-blind rewrite validation baseline
Experimental completeness	3/10	Seven of twelve main development runs are finalized; downstream locked/ablation studies remain absent
Statistical evidence	4/10	Paired development fusion evidence exists; locked architectural inference is absent
Human linguistic evidence	0/10	Dataset review and system ratings absent
Manuscript/submission package	4/10	Draft exists; core results, metadata, target adaptation absent

Overall blocking-gate rating: 5.5/10. This is not an average of cosmetic assets; it reflects that the primary hypothesis has no result.

Mandatory work before submission

1. Finish all 12 M0–M3 main runs and the registered optimizer, bidirectional-ID, external-baseline, factorial, confirmatory, and optimization studies.
2. Freeze the protocol before any main neural test access, retaining the prior classical/smoke disclosure.
3. Run all 12 locked main evaluations, at least 30 locked ablation evaluations, external baselines, robustness, routing, and registered statistics.
4. Require the primary M3–M2 source-OOD macro-chrF++ 95% hierarchical paired interval to exclude zero and its registered randomization test to survive familywise correction.
5. Demonstrate that M3 and the fixed M3 fused view are competitive with the fair inferred-dialect rewrite and credible external baselines—not only M2— while retaining the gold-dialect rewrite solely as an oracle; inspect MYM/NOA/CHI error modes.
6. Report IID, source-OOD, RAJ zero-shot, romanized, worst-dialect, calibration, throughput, memory, and retention trade-offs together.
7. Complete all 230 native dataset-review rows and blinded multi-rater system evaluation, including adequacy, fluency, hallucination, agreement, qualifications, consent, compensation, and ethics/IRB determination.
8. Select an actual target journal, verify its current quartile and policies, complete author metadata/declarations, expand the manuscript, format full references, and curate figures.

Final decision

Do not submit now. The project has a credible Q1-level research design, but the available evidence is an engineering/development package, not a completed paper. Q1 competitiveness becomes plausible only if the locked M3 result is positive, statistically robust, competitive with strong baselines, and confirmed by native-speaker evaluation. If M3 remains near the present M0/M1 range, the normalization system will need substantive modeling or decoding improvement before journal submission.

Exact evidence sources

- Dataset: `reports/FINAL_DATASET_REPORT.md` and

- tables/paper/dataset_counts.csv
- Tokenizer: reports/tokenizer/TOKENIZER_FREEZE_REPORT.json, reports/tokenizer/tokenizer_intrinsic_screen.csv, and reports/tokenizer/tokenizer_proxy_summary.csv
- Foundation/pilots: runs/*/1701/training_report.json and tables/paper/*.csv
- Main validation: runs/task/boichitro_q1_v1/*/*/stage_*/best_selection.json
- Consolidated main rows: reports/model/main_validation_results_current.csv with run-count/hash audit in reports/model/development_results_snapshot.json
- Dialect/class summaries: reports/model/main_validation_by_dialect_all_current.csv and reports/model/main_identification_by_class_all_current.csv
- Classical baselines: reports/model/classical_*_results.csv and metrics/
- Source-blind repair: reports/model/source_blind_normalization_baseline_audit.json, reports/model/cross_task_input_firewall.json, reports/model/normalization_fusion_selection_v2.json, reports/model/id_fusion_selection.json, and reports/model/development_fusion_uncertainty.json
- Fixed no-retuning transfer: reports/model/development_fusion_architecture_transfer.json
- Systems: reports/model/gb10_model_benchmark.json
- Disclosure: reports/PRIOR_TEST_ACCESS_DISCLOSURE.md
- Readiness: reports/Q1_JOURNAL_READINESS_AUDIT.md
- Figure validation: figures/q1/VALIDATION_REPORT.md